

## Prefix Sum

## Outline

- Reduction operations
  - Parallel reduction
- Prefix sum
  - Parallel scan
  - Work-efficient scan
- Applications of scan
  - Stream compaction

Monday, March 19, 2012

Minglun Gong

2

## Reduction

- A class of operations involves:
  - A ordered set  $S = \{a_0, a_1, a_2, \dots, a_{n-1}\}$  of  $n$  numbers
  - A binary associative operator
- Examples of reduction operations:
  - Sum =  $\text{Reduce}(+, S) = a_0 + a_1 + a_2 + \dots + a_{n-1}$
  - Product =  $\text{Reduce}(\times, S) = a_0 \times a_1 \times a_2 \times \dots \times a_{n-1}$
  - Min =  $\text{Reduce}(\min, S) = \min(a_0, a_1, a_2, \dots, a_{n-1})$
- The output is a single number
  - Require  $O(N)$  time to computer on a sequential computer


Monday, March 19, 2012

Minglun Gong

3

## Parallel Reduction

- Technique for performing reduction on parallel computers
  - Compute the sum of 2 numbers in each step
  - Reduce the numbers in the set by half
- Require  $\log_2 N$  steps on a  $N$ -processor computer
  - Require  $N/M + \log_2 M$  on  $M$  processors ( $M < N$ )



Monday, March 19, 2012

Minglun Gong

4

## Speedup & Efficiency

- Speedup is the time it takes to complete an algorithm on 1 processor divided by the time it takes on  $N$  processors
  - Measures the gain of parallelizing an algorithm
- Speedup of parallel reduction is  $N/\log_2 N$ 
  - With  $M$  processors  $M < N$ , the speedup is  $N/(N/M + \log_2 M)$
- Efficiency is defined as the speedup divided by the number of processors used
  - Measures how well the processors are unitized
- Efficiency of parallel reduction is  $1/\log_2 N$ 
  - With  $M$  processors  $M < N$ , the speedup is  $1/(N/M + \log_2 M)$


Monday, March 19, 2012

Minglun Gong

5

## Parallel Reduction on GPU

- Put data in a square texture & perform 2D reduction
  - Render a quarter-sized texture each pass
- Use shader:
  - Fetch nearby 4 values & calculate the sum
- Use build-in texture sampling functionality:
  - Set sampler to linear
  - Fetch the value at the center of the 4 pixels



Monday, March 19, 2012

Minglun Gong

6

## Prefix Sum (a.k.a. Scan)

- Given a list of  $n$  numbers, compute the partial sums using only numbers on the left sides
  - Input:  $a_0, a_1, a_2, \dots, a_{n-1}$
  - Output:  $a_0, a_0+a_1, a_0+a_1+a_2, \dots, a_0+a_1+a_2+\dots+a_{n-1}$
  - Require  $O(N)$  on a sequential computer
- Two variants of scan:
  - Inclusive scan: add all numbers on the left and the number itself
  - Exclusive scan: only add numbers on the left
    - The first output is zero
    - The last number in the input list is not used


Monday, March 19, 2012

Minglun Gong

7

## Parallel Scan

- A commonly used building block for parallel algorithms
  - Require  $\log_2 N$  steps on a  $N$ -processor computer
- In each step  $k$ ,  $k$  from 0 to  $\log_2 N$ :
  - if  $i > 2^k$ , add number  $a[i]$  with  $a[i-2^k]$



Monday, March 19, 2012

Minglun Gong

8

## Algorithm Complexity

- On a computer with  $N$  processors:
  - The total time needed to complete is  $\log_2 N$
  - The speedup is  $N/\log_2 N$
  - The efficiency is  $1/\log_2 N$
- On a computer with  $M$  processors ( $M < N$ ):
  - The total number of addition operations needed is  $N \times \log_2 N$
  - The total time needed to complete the additions is  $(N \times \log_2 N)/M$
  - Reduce the redundant add operations can further improve processing speed

Monday, March 19, 2012

Minglun Gong

9

## Work Efficient Parallel Scan

- Based on the balanced tree data structure
  - Build a balanced binary tree on the input data, then traverse the tree to and from the root
  - Perform one add per tree node, resulting a total of  $O(N)$  addition operations
- The algorithm consists of 2 phases
  - Upsweep phase traverses the tree from leaves to root computing partial sums
  - Down-sweep phase traverses from the root to leaves, using the partial sums to build the scan


Monday, March 19, 2012

Minglun Gong

10

## Upsweep Phase

- // same operation as parallel reduction
- for ( $d = 1$  to  $\log_2 n$ ) {
  - for ( $i = 1$  to  $n/2^d - 1$ )
  - do in parallel {
    - $a_d[i] = a_{d-1}[2i] + a_{d-1}[2i+1]$




Monday, March 19, 2012

Minglun Gong

11

## Down-sweep Phase

- for ( $d = (\log_2 n) - 1$  downto 0) {
  - for ( $i = 0$  to  $n/2^d - 1$ )
  - do in parallel {
    - if ( $i > 0$ ) {
      - if ( $(i \bmod 2) \neq 0$ ) {
        - $a_d[i] = a_{d+1}[i/2]$
      - else {
        - $a_d[i] += a_{d+1}[(i/2) - 1]$



Monday, March 19, 2012

Minglun Gong

12

## Applications of Scan

- Radix sort
- Quicksort
- String comparison
- Lexical analysis
- Stream compaction
- Sparse matrices
- Polynomial evaluation
- Solving recurrences
- Tree operations
- Histograms

Monday, March 19, 2012

Minglun Gong

13

## Stream Compaction

- Generate a compact stream by removing unwanted items from the original stream
  - Input: an ordered set  $S$  & a predicate  $p$
  - Output: only elements  $v$  for which  $p(v)$  is true, preserving the ordering of the input elements
- Applications:
  - An important operation in collision detection & sparse matrix compression
  - Can be used to transform a heterogeneous vector, with elements of many types, into homogeneous vectors, in which each element has the same type


Monday, March 19, 2012

Minglun Gong

14

## Stream Compaction Example

- Remove  $\leq 4$  numbers from the input stream
- Create a bit stream
  - Label  $>4$  with 1
  - Label  $\leq 4$  with 0
- Apply exclusive prefix sum on the bit stream
- Store numbers into the addresses specified by the result of prefix sum
  - Require scatter support



Monday, March 19, 2012

Minglun Gong

15