

Parallel Prefix Computation

Advanced Algorithms & Data Structures

Lecture Theme 14

Prof. Dr. Th. Ottmann

Summer Semester 2006

Overview

- A simple parallel algorithm for computing parallel prefix.
- A parallel merging algorithm

Definition of prefix computation

- We are given an ordered set A of n elements and a binary associative operator \oplus .

$$A = \{a_0, a_1, a_2, \dots, a_{n-1}\}$$

- We have to compute the ordered set

$$\{a_0, (a_0 \oplus a_1), \dots, (a_0 \oplus a_1 \oplus \dots a_{n-1})\}$$

An example of prefix computation

- For example, if \oplus is $+$ and the input is the ordered set


$\{5, 3, -6, 2, 7, 10, -2, 8\}$

then the output is

$\{5, 8, 2, 4, 11, 21, 19, 27\}$

- Prefix sum can be computed in $O(n)$ time sequentially.

Using a binary tree



First Pass

- For every internal node of the tree, compute the sum of all the leaves in its subtree in a **bottom-up** fashion.

$$sum[v] := sum[L[v]] + sum[R[v]]$$

Parallel prefix computation

for $d = 0$ to $\log n - 1$ do

for $i = 0$ to $n - 1$ by 2^{d+1} do in parallel


$$a[i + 2^{d+1} - 1] := a[i + 2^d - 1] + a[i + 2^{d+1} - 1]$$

- In our example, $n = 8$, hence the outer loop iterates 3 times, $d = 0, 1, 2$.

When $d = 0$

- $d = 0$: In this case, the increments of 2^{d+1} will be in terms of 2 elements.
- for $i = 0$,
$$a[0 + 2^{0+1} - 1] := a[0 + 2^0 - 1] + a[0 + 2^{0+1} - 1]$$
or, $a[1] := a[0] + a[1]$

Using a binary tree



First Pass


- For every internal node of the tree, compute the sum of all the leaves in its subtree in a **bottom-up** fashion.

$$sum[v] := sum[L[v]] + sum[R[v]]$$

When $d = 1$

- $d = 1$: In this case, the increments of 2^{d+1} will be in terms of 4 elements.
- for $i = 0$,
$$a[0 + 2^{1+1} - 1] := a[0 + 2^1 - 1] + a[0 + 2^{1+1} - 1]$$
or, $a[3] := a[1] + a[3]$
- for $i = 4$,
$$a[4 + 2^{1+1} - 1] := a[4 + 2^1 - 1] + a[4 + 2^{1+1} - 1]$$
or, $a[7] := a[5] + a[7]$

The First Pass



- **blue**: no change from last iteration.
- **magenta**: changed in the current iteration.

The Second Pass


Second Pass

- The idea in the second pass is to do a topdown computation to generate all the prefix sums.
- We use the notation *pre*[*v*] to denote the prefix sum at every node.

Computation in the second phase

- $pre[root] := 0$, the identity element for the \oplus operation, since we are considering the $+$ operation.
- If the operation is max , the identity element will be $-\infty$.


Second phase (continued)



$$pre[L[v]] := pre[v]$$

$$pre[R[v]] := sum[L[v]] + pre[v]$$

Example of second phase



$$\begin{aligned}pre[L[v]] &:= pre[v] \\pre[R[v]] &:= sum[L[v]] + pre[v]\end{aligned}$$

Parallel prefix computation

for $d = (\log n - 1)$ downto 0 do

for $i = 0$ to $n - 1$ by 2^{d+1} do in parallel

$temp := a[i + 2^d - 1]$

$a[i + 2^d - 1] := a[i + 2^{d+1} - 1]$ (left child)

$a[i + 2^{d+1} - 1] := temp + a[i + 2^{d+1} - 1]$ (right child)

$a[7]$ is set to 0

Parallel prefix computation

- We consider the case $d = 2$ and $i = 0$


$$temp := a[0 + 2^2 - 1] := a[3]$$

$$a[0 + 2^2 - 1] := a[0 + 2^{2+1} - 1] \text{ or, } a[3] := a[7]$$

$$a[0 + 2^{2+1} - 1] := temp + a[0 + 2^{2+1} - 1] \text{ or,}$$

$$a[7] := a[3] + a[7]$$

Parallel prefix computation



- **blue**: no change from last iteration.
- **magenta**: left child.
- **brown**: right child.

Parallel prefix computation

- All the prefix sums except the last one are now in the leaves of the tree from left to right.
- The prefix sums have to be shifted one position to the left. Also, the last prefix sum (the sum of all the elements) should be inserted at the last leaf.
- The complexity is $O(\log n)$ time and $O(n)$ processors.

Exercise: Reduce the processor complexity to $O(n / \log n)$.

Proof of correctness

- Vertex x precedes vertex y if x appears before y in the preorder (depth first) traversal of the tree.

Lemma: After the second pass, each vertex of the tree contains the sum of all the leaf values that precede it.


Proof: The proof is inductive starting from the root.

Proof of correctness

Inductive hypothesis: If a parent has the correct sum, both children must have the correct sum.

Base case: This is true for the root since the root does not have any node preceding it.

Proof of correctness



- **Left child:** The left child $L[v]$ of vertex v has exactly the same leaves preceding it as the vertex itself.

Proof of correctness

- These are the leaves in the region A for vertex $L[v]$.
- Hence for $L[v]$, we can copy $pre(v)$ as the parent's prefix sum is correct from the inductive hypothesis.

Proof of correctness

- **Right child:** The right child of v has two sets of leaves preceding it.
 - The leaves preceding the parent (region A) for $R[v]$
 - The leaves preceding $L[v]$ (region B).

$pre(v)$ is correct from the inductive hypothesis.

Hence, $pre(R[v]) := pre(v) + sum(L[v])$.